


Spark + Hive

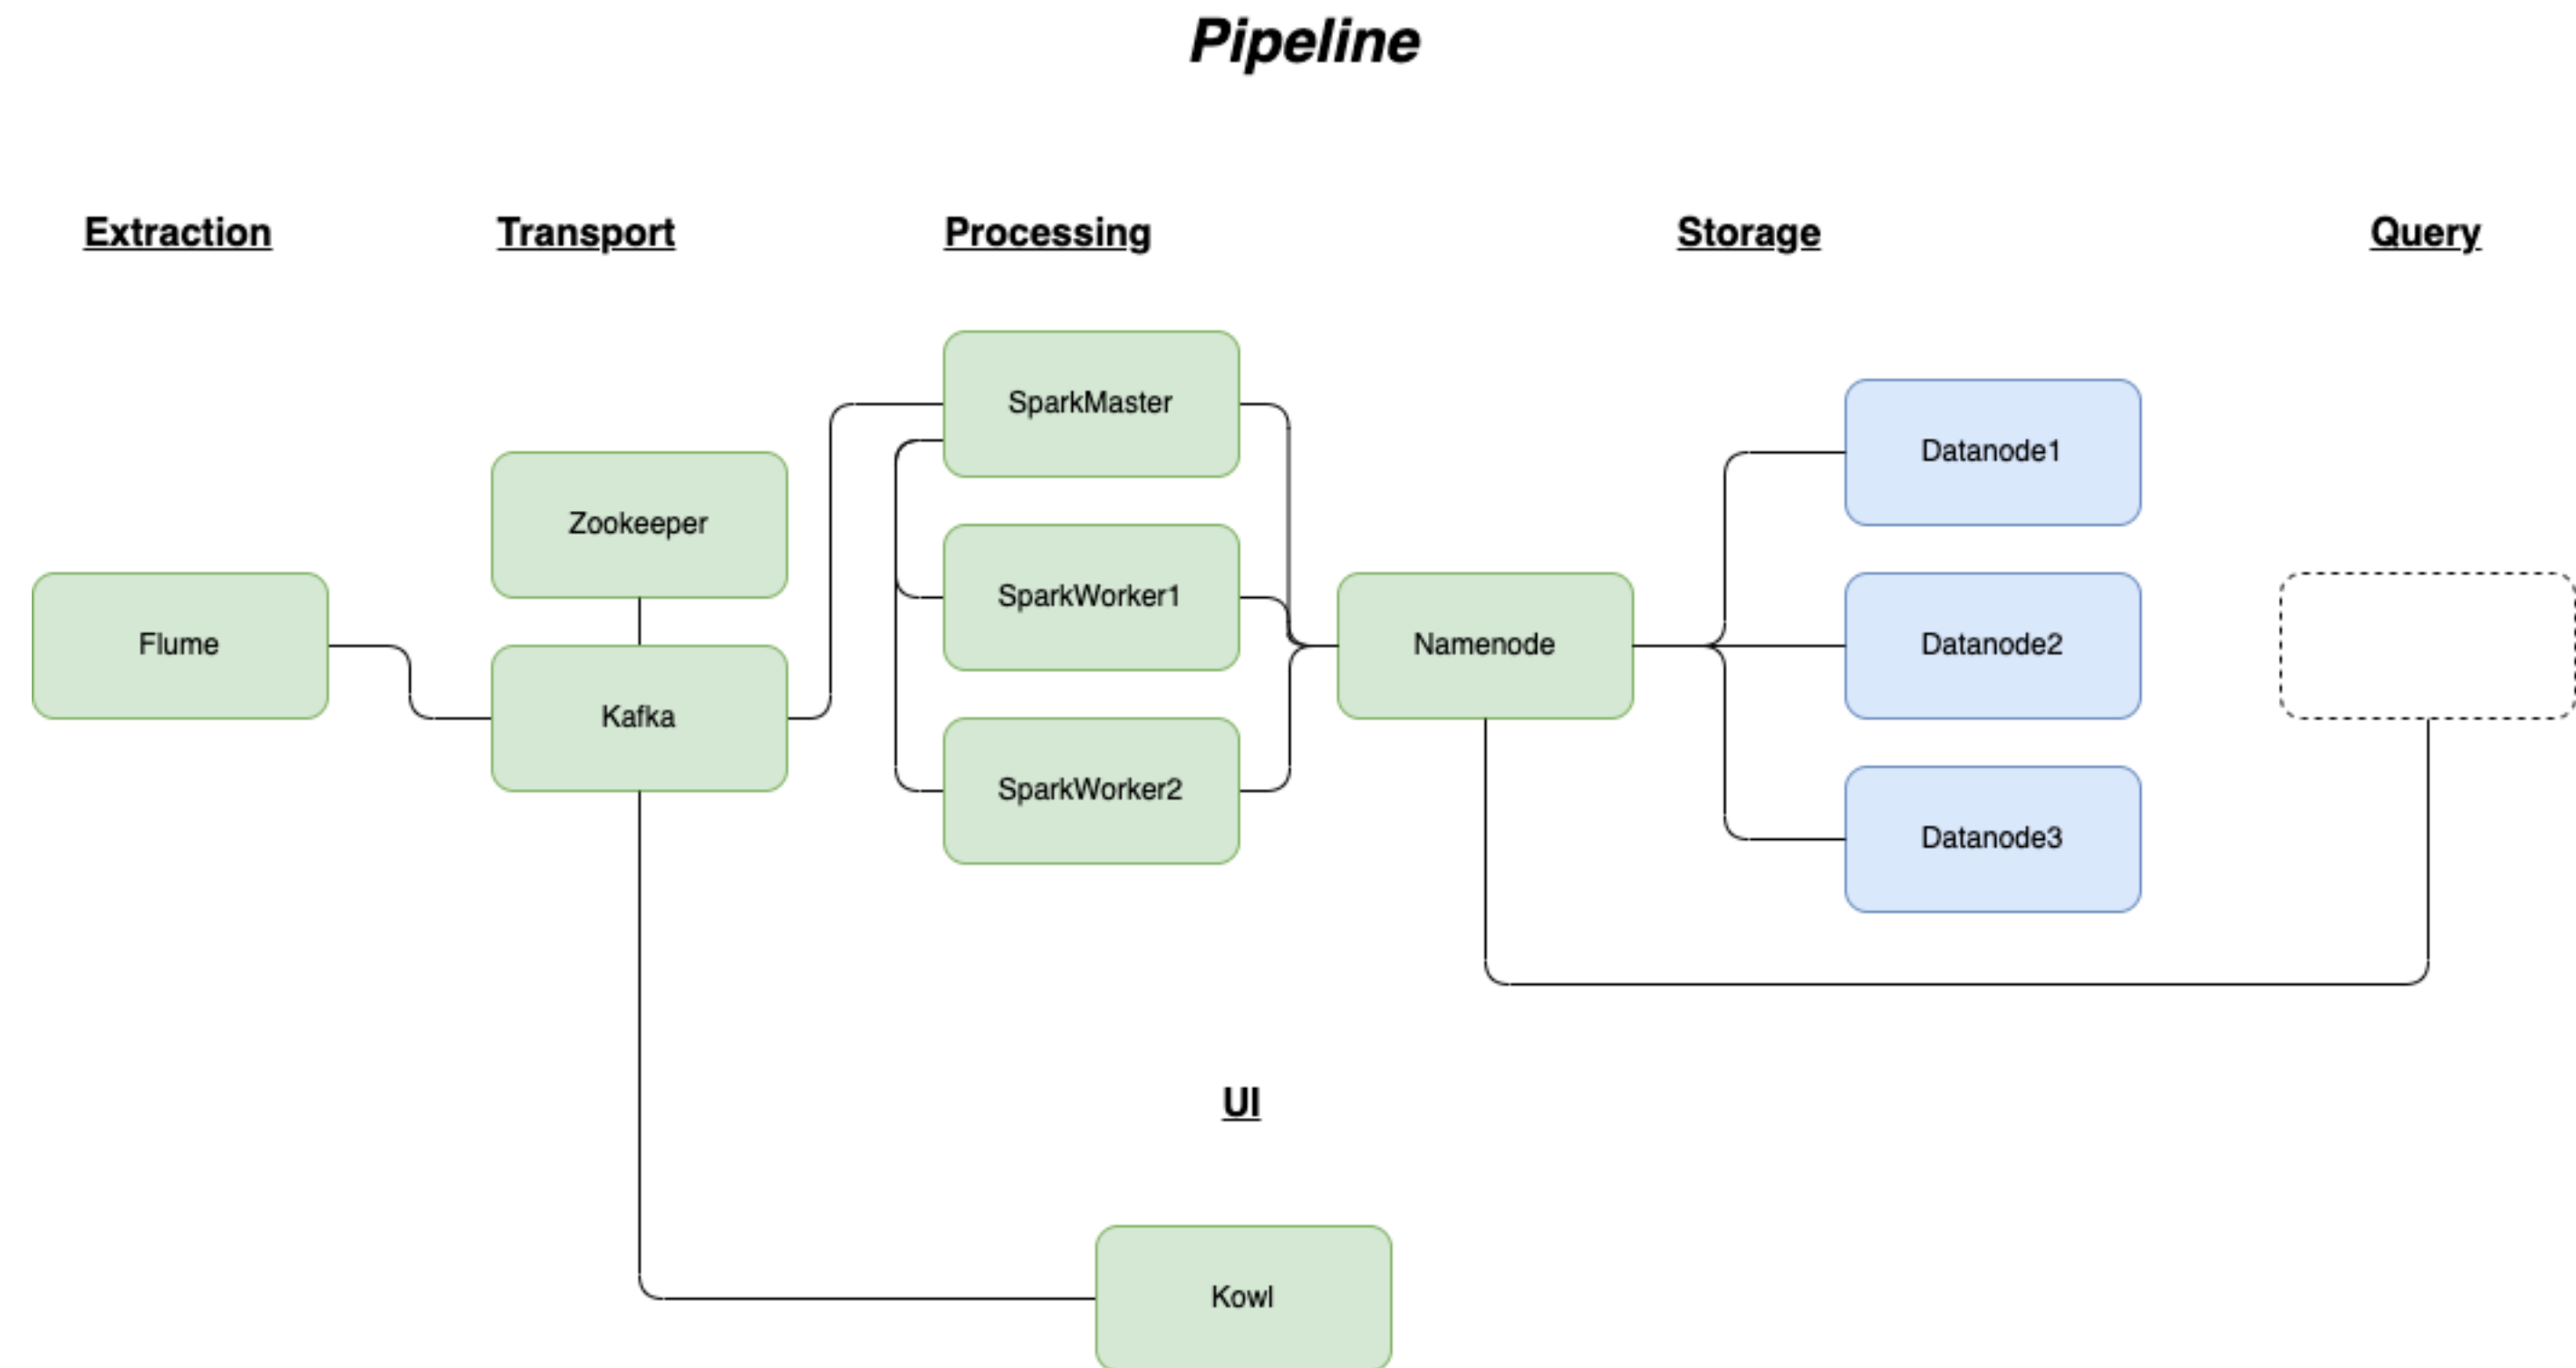


Big Data E22

Context

What have we been doing?

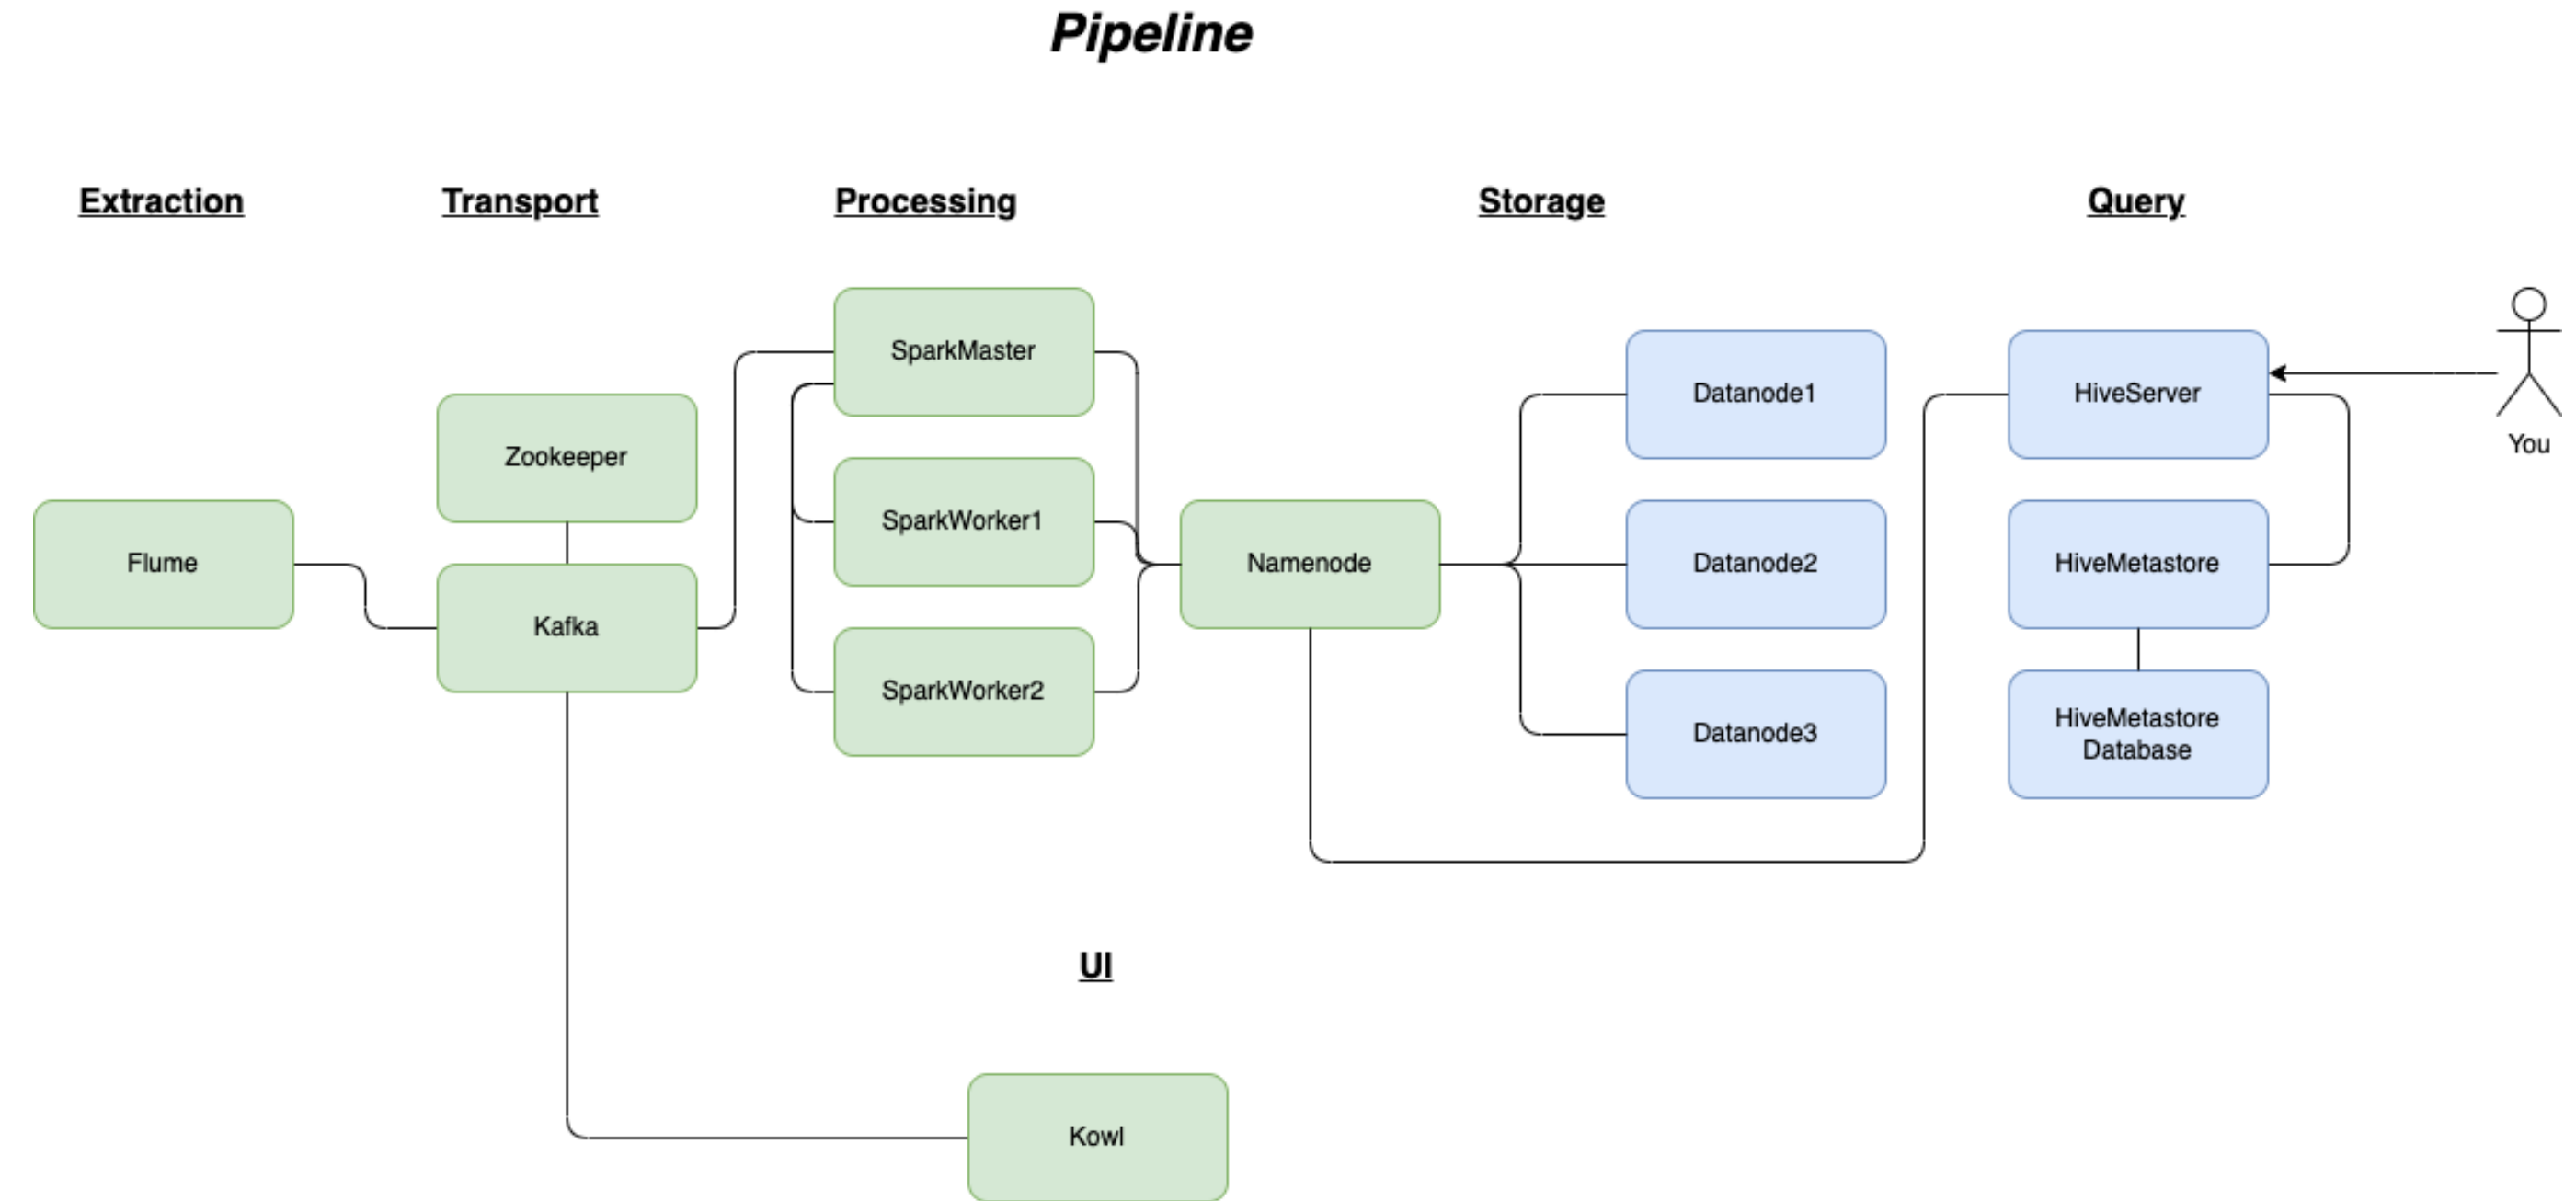
- A data-pipeline to ingest, process and store data.
- Last time we introduced distributed processing.



Context

What are we doing today?

- Revisit Spark Streaming
- Adding hive to our existing cluster
 - Query the wonderland!



Exercise 1

Composing a cluster with Hive

- To work with Hive we need to setup a cluster with a Hive server and a Hiver metastore.
- To do this you need to do the following.
 1. Teardown previous clusters e.g. Hadoop and Kafka. You can use `docker system prune` after stopping the stacks.
 2. `cd ./lecture05-exercises/`
 3. Examine the `docker-compose.yml` file
 4. Run `docker compose up -d`. Remove the argument `-d` if you need to see what happens.

Exercise 2

Upload Alice in Wonderland to HDFS (if you haven't already)

- `docker exec -ti namenode bash`
- `apt update`
- `apt install wget`
- `wget -O alice-in-wonderland.txt https://www.gutenberg.org/files/11/11-0.txt`
- `hdfs dfs -mkdir /txt`
- `hdfs dfs -put alice-in-wonderland.txt /txt/`

Exercise 3

Count words in Alice in Wonderland with Hive!

- Exec into the hive server (beeline is the name of the Hive CLI), and connect to the CLI to the server:
 1. `docker exec -ti hive-server beeline`
 2. `!connect jdbc:hive2://localhost:10000`
 3. `username and password is "hive"`
- Now lets show existing tables, create a new table, and load some data into it!
 4. `SHOW TABLES;`
 5. `CREATE TABLE lines (line STRING);`
 6. `LOAD DATA INPATH 'hdfs://namenode:9000/txt/alice-in-wonderland.txt' OVERWRITE INTO TABLE lines;`
- Awesome! Now we want structure our new data, so lets create another table for word counts.
 7. `CREATE TABLE word_counts`
`AS SELECT word, count(1)`
`AS count FROM (SELECT explode(split(line, ' ')) AS word FROM lines) w`
`GROUP BY word`
`ORDER BY word;`
- Finally we should query the word counts and see how they look!
 8. `SELECT * FROM word_counts ORDER BY count DESC LIMIT 10;`

Exercise 4

Check out HDFS

- Exec into the namenode.
- List files in the folder you put the `alice-in-wonderland.txt` in.
- Is the file there? If yes why? If no why not?
 - Hive is hungry beast 🤪 What are the benefits and disadvantages of the behaviour you discovered?

Exercise 5 Part 1

Hive external tables

- Why should we use Hive external tables?
 - Loose coupling with the data.
 - Data can be managed by more than Hive.
 - To avoid that dropping tables in Hive deletes data.
- Let's start out by cleaning up our Hive tables!
 - `DROP TABLE word_counts, lines`
 - `SHOW TABLES;` - verify the tables are gone 🙄
- Upload `alice-in-wonderland.txt` to HDFS again. You can follow Exercise 1 if in doubt on how to do that.

Exercise 5 Part 2

Hive external tables

- Lets create an external table!
 - `CREATE EXTERNAL TABLE lines (line string) LOCATION 'hdfs://namenode:9000/txt';`
- Now lets recreate the `word_counts` table
 - See previous slide.
 - Verify that it is recreated with `SELECT * FROM word_counts ORDER BY count DESC LIMIT 10;`
- Now lets add another book into HDFS and query from them both!
 - `hdfs dfs -put alice-in-wonderland.txt /txt/alice-in-wonderland2.txt` on the namenode
 - `SELECT COUNT(*) FROM lines;`
 - `SELECT INPUT__FILE__NAME FROM lines GROUP BY INPUT__FILE__NAME;`
- What happened? What results did you see?

Exercise 6

Sentiment exercise with Hive!

- Next we want to query sentiment scores with Hive. To do this we should first add the sentiment files to HDFS. The files are located in [lecture05-exercises/sentiment-files](#), and should be added into a HDFS folder as usual. (See Exercise 2 if in doubt)
- Now we need to load these files into Hive:
 - Use the **LOAD DATA** query as in exercise 3.
 - The files should be loaded into two tables - **positive_words_raw** and **negative_words_raw**.
- Now we need to structure the data so it is easier to work with:
 - For both positive and negative words, you must now manipulate the data into two new tables consisting of a single column with a unique word in each row.
 - The tables should be called **positive_words** and **negative_words**.
- Finally we want to do a query where we join these tables with the **word_counts**.
 - Try different queries, and see if you can figure out how to join the positive and negative words with the word counts.